




A. The better labs optimize, the less their data reveal



B. Profile likelihood of σ^* when κ is free



C. Compute-optimal tokens per parameter M^*



D. Sampling distribution of $\hat{\sigma}^*$

